

CONCURRENCY AND REPLICA CONTROL :-

Handling large clients

Handling multiple servers

CONCURRENCY CONTROL :-

RPC → Remote Procedure call.

Abstraction for processes to call functions in other processes.

For object based settings

④ Code reuse

↳ RMI (Remote method invocation)

LOCAL PROCEDURE CALL → Call from one function to another within the same process.

④ Stack usage to pass arguments and return values.

⇒ Object access via pointers. (because same address space)

Follows Exactly-once semantics.

RPC : → Caller and callee are in different processes.

④ Object references via global pointers.

EX. Obj. Address = IP + Port + Obj. Number

Basically crosses process boundaries.

FAILED CALL :-

→ Request message dropped

→ Reply message dropped

→ Called process fails

④ These cases are hard to distinguish

before execution
after execution

MULTIPLE CALL :-

→ Request message duplicated by the network

Possible Semantics

AT MOST ONCE

AT LEAST ONCE

MAYBE

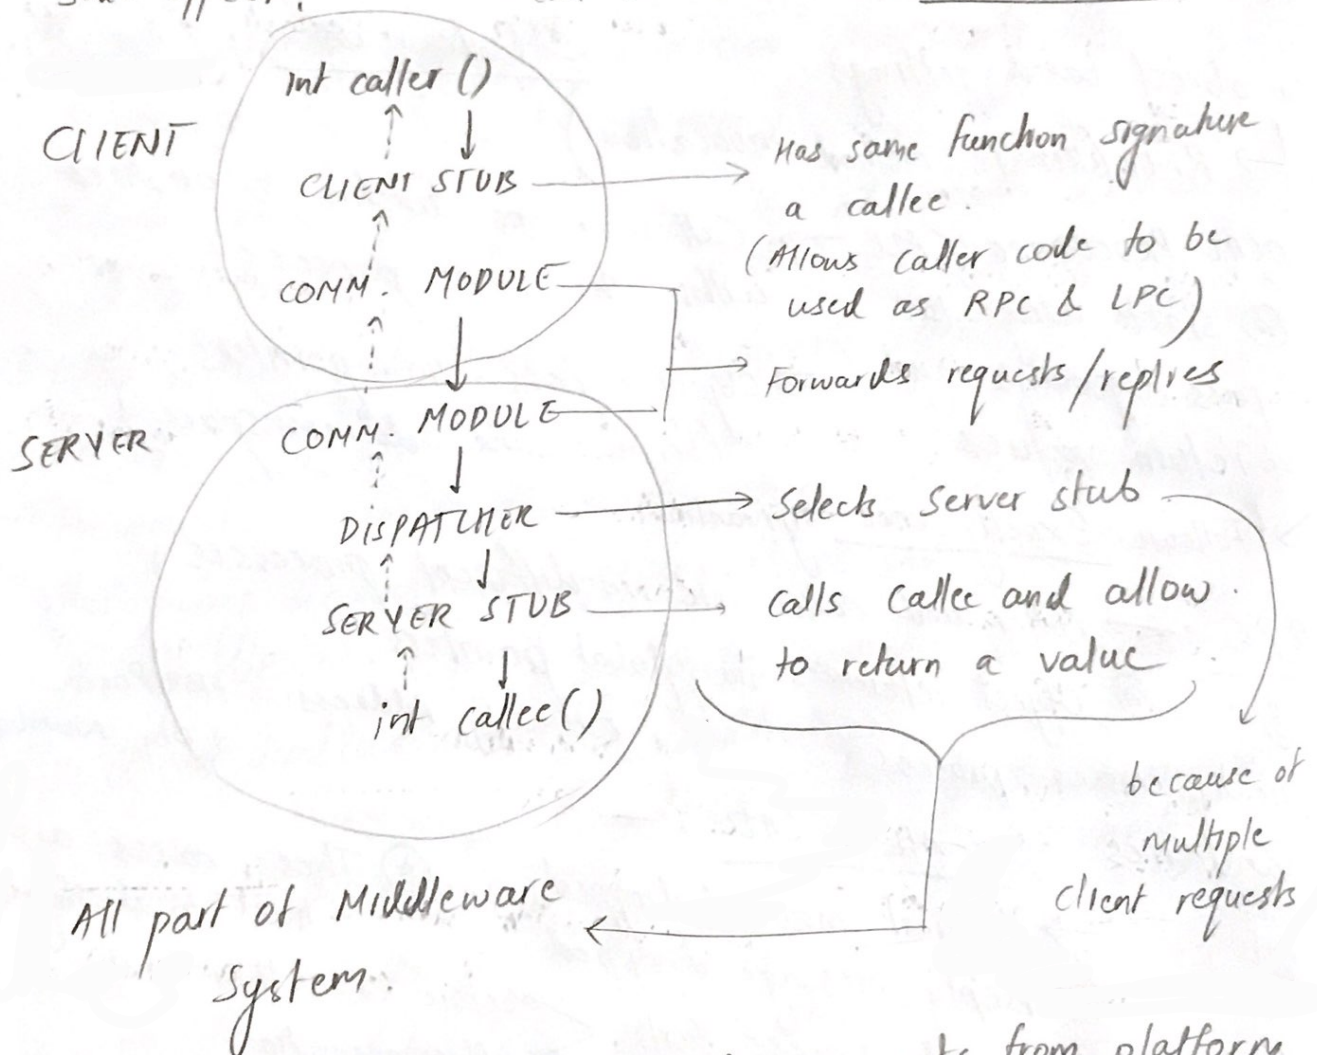
→ JAVA RMI

→ Sun RPC

→ CORBA

	<u>RETRANSMIT</u>	<u>FILTER DOPLICATES</u>	<u>RE-EXCUTE OR RETRANSMIT REPLY</u>
AT LEAST ONCE	Yes	NO	Re-execute
AT MOST ONCE	Yes	Yes	Retransmit
MAYBE	NO	N/A	N/A

② Idempotent operations can be repeated without any side effect. → can be used with AT LEAST ONCE



MARSHALLING → caller converts arguments from platform dependent format to COMMON DATA REPRESENTATION (CDR)

③ CDR is used by Middleware to become platform independent. REVERSE PROCESS → Unmarshalling.

TRANSACTIONS :

→ Series of operations executed by client
→ Each operation is an RPC to a server

either commits
all operations
at server

or Aborts
all operations
and has no
effect

PROBLEMS :—

- ① Lost update Problem
(Solved by Isolation)
- ② Inconsistent Retrieval

REQUIREMENTS :—

- Atomicity
- Isolation → Indivisible from point of view of other transactions (Basically free from interference)
- Consistency
- Durability

SERIAL EQUIVALENCE :—

An interleaving of transactions is serially equivalent iff there exists ordering of those transactions one at a time which gives same end-result (for all the objects and transactions)

Read → Affects client
Write → Affects server.

CONFLICTING OPERATIONS

Combined effect depends on order of execution
EX. Read(x), write(x)

- ① Serial Equivalence achieved when pairs of conflicting operations are in same order for all the objects

- ② Must abort one of the transactions when conflict in 2 transactions detected.

2 APPROACHES : ① PESSIMISTIC ② OPTIMISTIC

Prevent access
Use exclusive locks (Read-write locks)

check at commit time

PROBLEM: Reduces concurrency.

TWO PHASE LOCKS :- (PESSIMISTIC CONCURRENCY CONTROL)

→ Transaction cannot acquire any locks after it has started releasing locks.

PHASE ①: Growing Phase
PHASE ②: Shrinking phase

} Guarantees serial equivalence

DISADVANTAGE: Deadlocks

* expensive

* Might not even be a deadlock

* Wasted work

FIX: ① Timeout: Abort transaction

② Deadlock Detection:

- * Keeping track of wait-for graph
- * Find cycles
- * Abort one or more transactions to break cycle

③ Deadlock Prevention:

1. Allow read-only access to objects
2. Allow pre-emption of some transactions
3. Compound locks → lock all objects at the beginning

② OPTIMISTIC CONCURRENCY CONTROL

1. BASIC APPROACH: Check for serial equivalence, abort if not satisfied → leads to cascading abort

2. TIME STAMP ORDERING:

↓
Transaction id determines

its position in serialization order (Abort if rule violated)

* Read only if write by lower id in past

* Write only if read-write by lower ids in past

Transactions who read dirty data of aborted transaction needs to be aborted

③ MULTI-VERSION CONCURRENCY: —

→ Per-transaction version of object is maintained and marked as tentative versions (alongside committed version at server).

✓
Has timestamp

→ On read-write, mark correct version to read or write from
↳ based on transaction id

Eventual consistency → similar to Optimistic concurrency

RIAK KEY-VALUE STORE → Vector clock implements Causal ordering.

④ Sibling value resolved by user or application.

REPLICATION CONTROL :-

↳ Fault tolerance

↳ Load balancing

Higher availability? $1-f$ per server $\Rightarrow 1-f^k$ any 1 server is working

CHALLENGE: → Transparency (replication must be invisible) \Rightarrow Frontend

↳ Consistency (sees single consistent copy)

④ Concept of Replicated State Machines

① PASSIVE REPLICATION

→ Elected master in the system \Rightarrow Determines total ordering of all updates.
(on leader failure, run election)

provides replication transparency

② ACTIVE REPLICATION

Frontend multicasts the request to read/write to entire replica group
Total order on multicast \Rightarrow same inputs to replicas. \rightarrow can use any ordering.

B. Failures in active replication dealt by Virtual Synchrony

ONE-COPY SERIALIZABILITY :-

A concurrent execution of transaction in a replicated database is one-copy-serializable if it is equivalent to serial execution of these transactions on a single log. copy of database

FINAL OBJECTIVE = Serial equivalence + One-copy Serializability

- ① A transaction may touch different servers for different objects. Commit must commit to all or no servers (Atomic commit).

① ONE-PHASE COMMIT \Rightarrow special server Coordinator initiates atomic commit

PROBLEM :- ① Problems at a single server like corrupted object
② Server may crash before receiving commit msg.

② TWO-PHASE COMMIT

① Coordinator sends PREPARE message to servers. retrievable after crash.

② Server saves updates to disk and reply YES or NO

③ If all votes received are YES within the timeout, it sends commit message. Otherwise abort.

④ On commit, server commits and sends ACK. server can poll if no decision received

To deal with COORDINATOR crash, it logs all messages, decisions on disk. After recovery, new election.