

## CONCURRENCY AND REPLICA CONTROL :-

Handling large clients

Handling multiple servers

### CONCURRENCY CONTROL :-

RPC → Remote Procedure call.

For object based settings

Abstraction for processes to call functions in other processes.

↳ RMI (Remote method invocation).

LOCAL PROCEDURE CALL → Call from one function to another within the same process.

④ Stack usage to pass arguments and return values. ⇒ object access via pointers.  
(because same address space)

Follows Exactly-once semantics.

RPC :→ Caller and callee are in different processes.

↳ Object references via global pointers.  
Basically crosses process  
ex. obj. address = IP + Port + Obj. Number

boundaries.

### FATLED CALL :-

→ Request message dropped

→ Reply message dropped

→ Called process fails

④ These cases are hard to distinguish

before execution

after execution

### MULTIPLE CALL :-

→ Request message duplicated by the network

### Possible Semantics

AT MOST ONCE

AT LEAST ONCE

MAYBE  
→ CORBA

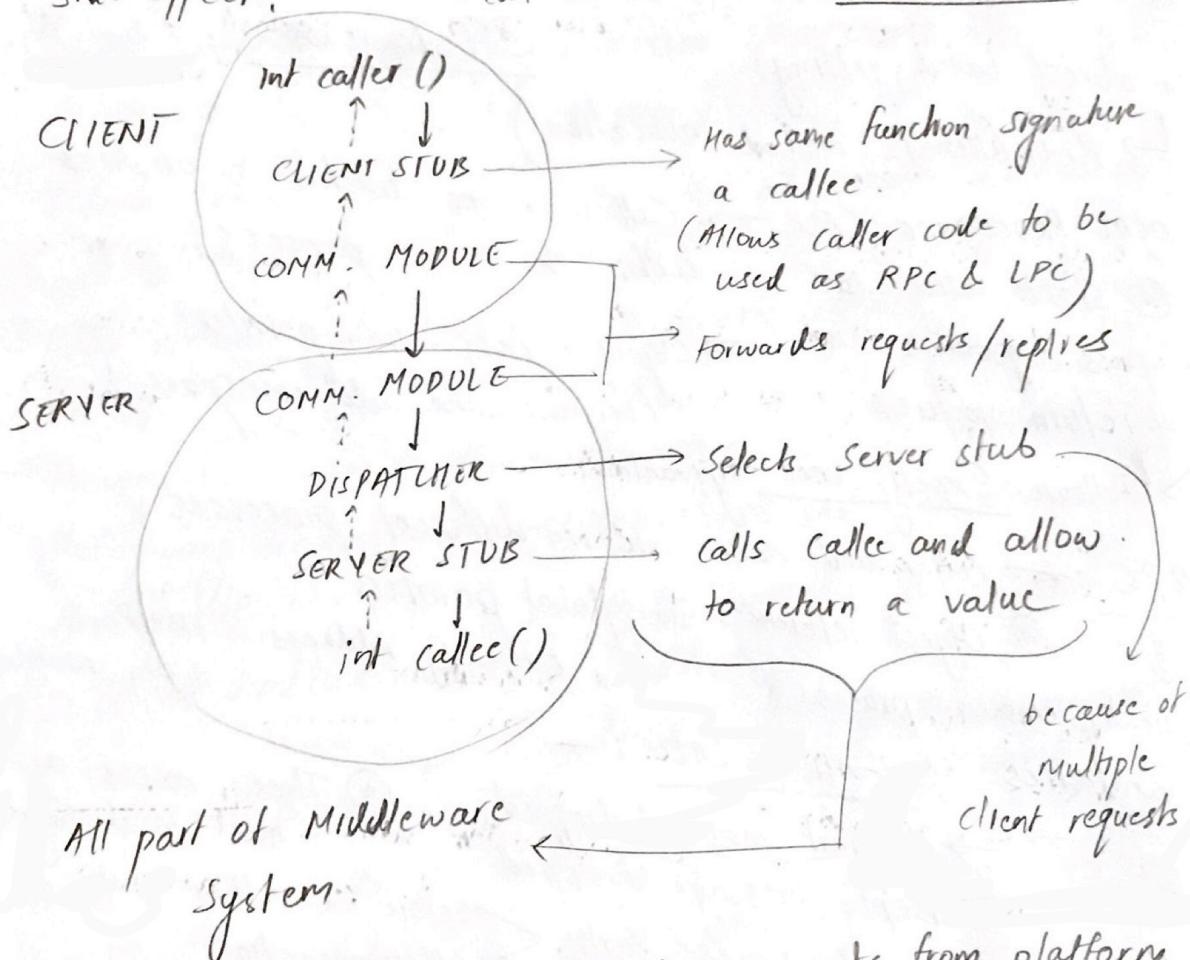
→ JAVA RMI

→ Sun RPC

	<u>RETRANSMIT</u>	<u>FILTER Duplicates</u>	<u>RE-EXECUTE OR RETRANSMIT REPLY</u>
AT LEAST ONCE	Yes	No	Re-execute
AT MOST ONCE	Yes	Yes	Retransmit
MAYBE	No	N/A	N/A

② Idempotent Operations can be repeated without any side effect.

→ can be used with AT LEAST ONCE



MARSHALLING → Caller converts arguments from platform dependent format to COMMON DATA REPRESENTATION (CDR)

③ CDR is used by Middleware to become Platform independent

REVERSE PROCESS → Unmarshalling.

TRANSACTIONS :-

- Series of operations executed by client
- Each operation is an RPC to a server

- either commits
- all operations at server
- or Aborts
- all operations and has no effect

PROBLEMS :-

- ① Lost update Problem  
(Solved by Isolation)
- ② Inconsistent Retrieval

REQUIREMENTS :-

- Atomicity
- Isolation → Indivisible from point of view of other transaction
- Consistency
- Durability (Basically free from interference)

SERIAL EQUIVALENCE :-

An interleaving  $\sigma$  of transactions is serially equivalent iff there exists ordering  $\sigma'$  of those transactions one at a time which gives same end-result (for all the objects and transactions)

- Read  $\rightarrow$  Affects client
- write  $\rightarrow$  Affects server.

- ③ Serial Equivalence achieved when pairs of conflicting operations are in same order for all the objects

CONFLICTING OPERATIONS

combined effect depends on order of execution  
ex. Read(x), write(x)

- ④ Must abort one of the transactions when conflict in 2 transactions detected

2 APPROACHES : ① PESSIMISTIC ② OPTIMISTIC

↓ Prevent access → check at commit time  
use exclusive locks (Read-write locks)

PROBLEM: Reduces concurrency

## TWO PHASE LOCKS :- (PESSIMISTIC CONCURRENCY CONTROL)

→ Transaction cannot acquire any locks after it has started releasing locks.

PHASE ①: Growing Phase → Guarantees serial equivalence

PHASE ②: Shrinking phase

### DISADVANTAGE: Deadlocks

FIX : ① Timeout : Abort transaction \* expensive  
\* Might not even be a deadlock  
\* Wasted work

② Deadlock Detection : \* Keeping track of wait graph  
\* Find cycles  
\* Abort one or more transactions to break cycle

③ Deadlock Prevention : 1. Allow read-only access to objects  
2. Allow pre-emption of some transactions  
3. Compound locks → lock all objects at the beginning

## (2) OPTIMISTIC CONCURRENCY CONTROL

1. BASIC APPROACH : Check for serial equivalence, abort if not satisfied → leads to cascading abort.

### 2. TIME STAMP ORDERING :

↓  
Transaction id determines

its position in serialization order (Abort if rule violated)

\* Read only if write by lower id in past

\* Write only if read/write by lower ids in past

Transactions who read dirty data of aborted transaction needs to be aborted

### ③ MULTI-VERSION CONCURRENCY : —

- Per-transaction version of object is maintained and marked as tentative versions (alongside committed version at server).  
Has timestamp
- On read/write, mark correct version to read or write from
  - ↳ based on transaction id
- Eventual consistency → similar to optimistic concurrency
- RIAK KEY-VALUE STORE → Vector clock implements  
④ Sibling value resolved by causal ordering.  
user or application.

### REPLICATION CONTROL :-

- Higher availability?  $1-f$  per server  $\Rightarrow 1-f^k$  any 1 server is working
- CHALLENGE: → Transparency (replication must be invisible)  $\Rightarrow$  Frontend
- ④ Concept of Replicated State Machines

### ① PASSIVE REPLICATION

- Elected master in the system  $\Rightarrow$  determines total ordering of all updates

provides replication transparency

- ### ② ACTIVE REPLICATION
- Frontend multicasts the request to read/write to entire replica group
  - Total order on multicast  $\Rightarrow$  same inputs to replicas  $\Rightarrow$  can use any ordering.

④ Failures in active replication dealt by Virtual Synchrony

### ONE-COPY SERIALIZABILITY :-

A concurrent execution of transaction in a replicated database is one-copy-serializable if it is equivalent to serial execution of these transactions on a single logical copy of database.

FINAL OBJECTIVE = Serial equivalence + One-copy Serializability

⑤ A transaction may touch different servers for different objects. Commit must commit to all or no servers (Atomic commit).

⑥ ONE-PHASE COMMIT  $\Rightarrow$  special server Coordinator initiates atomic commit

PROBLEM :- ① Problems at a single server like corrupted object  
② Server may crash before receiving commit msg.

⑦ TWO-PHASE COMMIT retrievable after crash.

① Coordinator sends PREPARE message to servers.

② Server saves updates to disk and reply YES or NO

③ If all votes received are YES within the timeout, it sends commit message. Otherwise abort.

④ On commit, server commits and sends ACK.  $\rightarrow$  server can poll if no decision received

To deal with COORDINATOR crash, it logs all messages, decisions on disk. After recovery, new election.