

CONSENSUS PROBLEM : —

DISTRIBUTED SYSTEMS

Reliable
Multicast

Membership/
Failure
Detection

Leader Election

Mutual
Exclusion

① COMMON : Coordinate and agree on a value of some parameter

CONSENSUS : —

Validity

Integrity

Non-triviality

Leader

Access
to critical resource

status of
failed process

Ordering of
messages

SYSTEM MODELS

SYNCHRONOUS

Time bound
on messages

Drift bound

Each
process
is bounded

Consensus Solution

Possible

* Use time slicing or
analyse different periods of
rounds based on above
bounds

→ Reverse induction on
faulty nodes

ASYNCHRONOUS

More general

No bound
on process
execution

No bound
on drift

No bound
on message
transmission
delays

Consensus Solution

Impossible

Popular Solutions

PAXOS / SIMPLY

PAXOS :

SAFETY → No 2 processes decide different values

EVENTUAL LIVENESS → Consensus reached eventually

① Move to next round.

PHASE-1 : Leader Election

- if timeout or message from next round
- ⇒ Potential leader chooses unique ballot id and send to all processes
 - ↳ Higher than anything seen
- ⇒ Processes wait and respond to highest ballot id
- ⇒ Process includes v in the response
- ⇒ If Quorum respond OK, then you are the leader
 - ↳ If failed quorum, start next round
- ② Potential leader withdraws if it sees higher value

PHASE-2 : Proposal (Bill)

- Leader sends proposed value v to all
 - use $v = v'$ if some process already decided in a previous round
- Recipients log on disk and respond OK

PHASE-3 : Decision (Law)

- If leader has majority of OK, it sends decision to all recipients

To ensure safety

Previous and new quorum intersect and hence same value used by new quorum.

Eventual liveness

Leader fails

Messages dropped

Process fails.